

Welcome

This is the source code for MpdFS described in our paper.
The system is tested on Ubuntu 12.04.2 64-bit with GCC 4.6.3.

Installation

These are the required libraries that users need to download separately.
Brackets denote the package names in Debian and Ubuntu platforms.
Users can use apt-get to install the required libraries.

- build-essential (build-essential)
- scons (scons)
- boost libraries version 1.5x (libboost-dev, libboost-program-options-dev, libboost-thread-dev, libboost-filesystem-dev)
- FUSE (libfuse-dev)
- openssl (libssl-dev)
- pkg-config (pkg-config)

For Debian and Ubuntu users:

```
sudo apt-get install build-essential scons libboost-dev libboost-program-options-dev libboost-thread-dev libboost-filesystem-dev libfuse-dev libssl-dev pkg-config
```

The following libraries have to be compiled and installed manually:

Google Protocol Buffers

```
$ cd lib/protobuf-2.4.1
$ ./configure; make clean; make
$ sudo make install
$ sudo ln -s /usr/local/lib/libprotobuf.so.7 /usr/lib/libprotobuf.so.7
$ sudo ln -s /usr/local/lib/libprotoc.so.7 /usr/lib/libprotoc.so.7
```

MongoDB C++ Driver (only required for MASTER)

```
$ cd lib/mongo
$ sudo scons install
```

Preparation

Compile

The program can be compiled using Linux make.
Running make on the MpdFS root will automatically compile all MpdFS components.

```
$ make
```

For optimal performance, turn off debug mode by editing the following flags in Makefile:

```
OPTIMIZE := -O3
EXTRA_CFLAGS := -std=c++0x -DDEBUG=0 -D_FILE_OFFSET_BITS=64
```

Setting up MongoDB

We suggest hosting the MongoDB on the MASTER for optimal performance.

- Set up MongoDB environment

On the MongoDB host, create the database parent directory

```
$ sudo mkdir -p /data/db
```

Change the MongoDB address "[mongodb_ip]:[mongodb_port]" in "mongo-mpdfs.js" if necessary

```
db = connect("localhost:27017/mpdfs");
```

Change the default MongoDB password in "mongo-mpdfs.js" for security reasons

```
"pwd": "cf8b7fa0b1bcd277da8217f04f498bd0"
```

Note

- "pwd" should be an MD5 hash, which can be generated by running `echo 'NEWPASSWORD' | md5sum` in Linux
- Change to the directory containing the MongoDB binary

```
$ cd mongo/bin
```

- Start MongoDB daemon

```
$ sudo ./mongod --fork --dbpath /data/db --logpath /var/log/mongodb.log --logappend
```

- Import MongoDB Setting

```
$ ./mongo ../../mongo-mpdfs.js
```

Adjusting XML Settings

- Copy one of the two set of example XML files from the directory to the MpdFS root (same level as the executables)
For standalone setup, `cp example_xml/standalone/* .`
For distributed setup, `cp example_xml/distributed/* .`
- example_xml/standalone/
 - Example XML files for a standalone setup which runs all components on localhost (127.0.0.1)
- example_xml/distributed/
 - Example XML files for a distributed setup
 - MONITOR on 192.168.0.11:53000

- MASTER on 192.168.0.12:50000
 - MongoDB daemon on 192.168.0.12:27017
- Adjust the following essential parameters in the XML files according to the cluster setup:
- `masterconfig.xml`
 - MASTER listen port
 - MongoDB IP address, listen port, password
- `csconfig.xml`
 - Update scheme (FO=0, FL=1, PL=2, PLR=3, MPLR=4) and reserved space size
- `monitorconfig.xml`
 - MONITOR listen port
- `clientconfig.xml`
 - Coding scheme (e.g., RS, RDP) and coding settings (e.g. n, k)
- `common.xml`
 - MASTER IP address and listen port
 - MONITOR IP address and listen port
- Refer to comments in the XML files for other individual settings (e.g., Chunkserver capacity)

Running MpdFS

Server-side

Copy the corresponding executable (MONITOR, MASTER or Chunkserver) and all the XML files to each server

Start the components in the following order

- MONITOR

```
$ ./MONITOR
```

- MASTER

```
$ ./MASTER
```

- Start Chunkserver one-by-one

```
$ ./CS [component_id] [network_interface]
```

Note

- Assign a `component_id` that is unique from MASTER, MONITOR, other Chunkservers and clients
- Use "ifconfig" to check the name of the network interface (e.g., eth0). If multiple interfaces exist, select the one that can be reached by other components

Client-side

On the client machine, create two directories

```
$ mkdir fusedir mountdir
```

Mount the FUSE volume with the following options

```
$ ./CLIENT_FUSE -o big_writes,large_read,noatime -f mountdir [component_id]
```

You can now access MpdFS through `mountdir`. The following command copies a file into the cluster and outputs the time spent:

```
$ time cp testfile mountdir/
```

Note

- Assign a `component_id` that is unique from MASTER, MONITOR, Chunkservers and other clients

Benchmarks

Test seq. read/write throughput using dd

After mounting the FUSE volume, test seq. write throughput by:

```
$ dd if=/dev/zero of=mountdir/ddtest count=256 bs=16M conv=fdatasync
```

After random writes(using IOzone), test seq. read throughput by:

```
$ dd if=mountdir/ddtest of=/dev/null count=256 bs=16M
```

Test random write throughput using IOzone

Download latest IOzone tarball from <http://www.iozone.org/>.

Compile IOzone from source

```
$ tar xf iozone3_424.tar
$ cd iozone3_424/src/current
$ make linux-AMD64
```

Copy the IOzone executable to Mpdroot, and execute the following:

```
$ ./iozone -Ra -+n -i0 -i2 -s 4g -r 128k -e -c -w -f mountdir/iozonetest
```

This will create a 4GB file and perform random read/write with 128KB record size.

Recovery

Automatic Recovery

MpdFS does not automatically recover by default.

To enable this features, uncomment the line `#define TRIGGER_RECOVERY` in `src/common/define.hh`

Note

- Adjust the sensitivity of recovery trigger by changing `SleepPeriod`, `DeadPeriod` and `updatePeriod` in `monitorconfig.xml`

Manual Recovery

Manual recovery is useful in benchmarks.

To simulate failure, kill the Chunkserver process on one of the Chunkserver hosts

```
$ kill -9 cs
```

To force failure detection and recovery, send a signal to the MONITOR process on the MONITOR host

```
$ kill -USR1 MONITOR
```

Traces

NFS Traces

MpdFS use Harvard NFS traces to evaluate (Harvard SOS Traces: <http://iotta.snia.org/traces/nfs/3378>)

Classifier

MpdFS-Classifier is built in the environment of python3

How to use:

1. run `readTraces`, read core records from old traces (from NFS traces)
2. run `createUser` and `createScheme`, create certain features needed for training
3. run `createFiles`, generate all required feature sets from the processed dataset and created rule set
4. run `createTrainSets`, generate dataset and training set
5. run `mainTest`, evaluate prediction performance
6. run `readOnlyMain`, predict read-only or not

Note

- Some file path parameters in the code need to be entered manually, for example "old_txt_path", "csv_path"...

Contact

Xiaosong Su (Email: sxsilly@163.com)